

# Machine Learning/Deep Learning stack using Singularity and Jupyter notebooks

*Mahidhar Tatineni*  
*UC Berkeley/LBL*  
*May 2, 2018*



# Machine Learning/Deep Learning on Comet via Singularity

- Bulk of the Singularity usage on Comet is for machine learning/deep learning applications.
- Lot of these packages are constantly upgraded and the dependency list is difficult to update in the standard Comet environment.
- Install options
  - Singularity image provides dependencies and user can compile actual application from source. e.g. Torch
  - Entire dependency stack and the application is in the image. e.g Keras with backend to TensorFlow and some additional python libraries
- Run options
  - Most cases are run on single GPU nodes (4 GPUs at most)
  - Can access this via Jupyter notebooks
  - Multi-node options are possible but difficult to set up via singularity.

# Tensorflow via Singularity (Single Node)

```
#!/bin/bash
#SBATCH --job-name="TensorFlow"
#SBATCH --output="TensorFlow.%j.%N.out"
#SBATCH --partition=gpu-shared
#SBATCH --nodes=1
#SBATCH --ntasks-per-node=6
#SBATCH --gres=gpu:1
#SBATCH -t 01:00:00

#Run the job
#

module load singularity
singularity exec /share/apps/gpu/singularity/sdsc_ubuntu_gpu_tflow.img lsb_release -a
singularity exec /share/apps/gpu/singularity/sdsc_ubuntu_gpu_tflow.img python -m tensorflow.models.image.mnist.convolutional
```

# Tensorflow via Singularity

- **Change to the examples directory:**

```
cd /home/$USER/UCB2018/TensorFlow
```

- **Submit the job:**

```
sbatch TensorFlow.sb
```

# Tensorflow Example: Output

**Distributor ID: Ubuntu**

**Description: Ubuntu 16.04 LTS**

**Release: 16.04**

**Codename: xenial**

I tensorflow/stream\_executor/dso\_loader.cc:108] successfully opened CUDA library libcublas.so locally

I tensorflow/stream\_executor/dso\_loader.cc:108] successfully opened CUDA library libcudnn.so locally

I tensorflow/stream\_executor/dso\_loader.cc:108] successfully opened CUDA library libcufft.so locally

I tensorflow/stream\_executor/dso\_loader.cc:108] successfully opened CUDA library libcuda.so.1 locally

I tensorflow/stream\_executor/dso\_loader.cc:108] successfully opened CUDA library libcurand.so locally

I tensorflow/core/common\_runtime/gpu/gpu\_init.cc:102] Found device 0 with properties:

name: Tesla K80

major: 3 minor: 7 memoryClockRate (GHz) 0.8235

pciBusID 0000:85:00.0

Total memory: 11.17GiB

Free memory: 11.11GiB

I tensorflow/core/common\_runtime/gpu/gpu\_init.cc:126] DMA: 0

I tensorflow/core/common\_runtime/gpu/gpu\_init.cc:136] 0: Y

I tensorflow/core/common\_runtime/gpu/gpu\_device.cc:838] Creating TensorFlow device (/gpu:0) -> (device: 0, name: Tesla K80, pci bus id: 0000:85:00.0)

Extracting data/train-images-idx3-ubyte.gz

...

Step 8500 (epoch 9.89), 11.6 ms

Minibatch loss: 1.601, learning rate: 0.006302

Minibatch error: 0.0%

Validation error: 0.9%

Test error: 0.9%

# Accessing via Jupyter Notebook

[1] Get an interactive node:

```
srun --pty --nodes=1 --ntasks-per-node=24 -p compute --  
reservation=UCB2018Res -t 02:00:00 --wait 0 /bin/bash
```

[2] Load the singularity module and get an interactive shell

```
module load singularity
```

```
singularity shell /share/apps/gpu/singularity/sdsc_ubuntu_tf1.1_keras_R.img
```

[3] Launch the notebook

```
ipython notebook --no-browser --ip="*" &
```

This will give you an address which has localhost in it and a token. Something like:

```
http://localhost:8888/?token=389587c9d1b69f8f595e7d8bfdd83c9961ed26b8b3f  
1bb3e
```

You can replace localhost with **comet-XX-YY.sdsc.edu** and then paste it into your browser. That should get you into the running notebook. From there everything should be working as a regular notebook. Note: This token is your auth so don't email/send it around (I already stopped the above link).



# MNIST Example

- **Keras, deep learning library in python - backend to Tensorflow or Theano. Today we will use Tensorflow.**
- **Deep learning refers to neural networks with multiple hidden layers that can learn increasingly abstract representations of the input data. For example in a image classification case**
  - first layers might learn local edge pattern
  - each subsequent layer (or filter) gets more complex reps
  - eventually the last layer can classify images - car, tree etc.
- **Convolutional Neural Networks - images are the input data - reduce the parametric space for tuning and effectively handle the high dimensionality of raw images.**
- **References:**
  - <https://cambridgespark.com/content/tutorials/deep-learning-for-complete-beginners-recognising-handwritten-digits/index.html>

# MNIST Example via Jupyter Notebook

- **Step 1: Follow the instructions for spinning up the Jupyter notebook**
- **Step 2: From the browser window, open the following notebook:**
  - LabMNIST\_Final.ipynb



# Multi-Node runs via Singularity

- **Easy for cases with MPI backends - we already saw this in the first talk.**
- **ML/DL frameworks can be a bit more complicated with process launches on remote nodes (need to be done via image)**
- **Two examples:**
  - Tensorflow - processes are launched once => just need to wrap the launch tasks.
  - MXNET : Lot of remote commands => need to patch the script that does this part.

# Script snippet from multi-node TF

**#Start the parameter server**

```
cd /scratch/$USER/$SLURM_JOBID
```

```
cp $SLURM_SUBMIT_DIR/example.sh .
```

```
cp $SLURM_SUBMIT_DIR/example.py .
```

```
./example.sh "ps" 0 2>&1 > $SLURM_SUBMIT_DIR/ps.$SLURM_JOBID.log &
```

**#Run the worker processes remotely**

```
ssh $H2 $SLURM_SUBMIT_DIR/run_worker.sh $SLURM_SUBMIT_DIR
```

```
/scratch/$USER/$SLURM_JOBID 0 > $SLURM_SUBMIT_DIR/worker0_$SLURM_JOBID.log  
&
```

```
sleep 10s
```

```
ssh $H3 $SLURM_SUBMIT_DIR/run_worker.sh $SLURM_SUBMIT_DIR
```

```
/scratch/$USER/$SLURM_JOBID 1 >
```

```
$SLURM_SUBMIT_DIR/worker1_$SLURM_JOBID.log
```

**\*\*\* For the multinode case, you need to untar the tensorflow.tar file in your home directory.**

# MXNET

- **Remote launch controlled by:**
  - incubator-mxnet/dmlc-core/tracker/dmlc\_tracker/ssh.py
- **Change made:**

```
prog = get_env(pass_envs) + ' cd ' + working_dir + '  
/opt/singularity/bin/singularity exec  
/oasis/scratch/comet/username/temp_project/singularity/mxn  
et-gpu.img ' + (' '.join(args.command))
```

# MXNET interactive snippet

```
scontrol show hostname > hosts
```

```
for NODE in `cat hosts`; do ssh $NODE cp -r /home/username/mxnet/incubator-mxnet/example/image-classification /scratch/$USER/$SLURM_JOBID; done
```

```
cd /scratch/$USER/$SLURM_JOBID/image-classification
```

```
singularity shell
```

```
/oasis/scratch/comet/username/temp_project/singularity/mxnet-gpu.img
```

```
python /home/username/mxnet/incubator-mxnet/tools/launch.py -n 2 --launcher  
ssh -H hosts python train_mnist.py --gpus 0,1,2,3 --network lenet --kv-store  
dist_sync
```

```
for NODE in `cat hosts`; do ssh $NODE killall -s 9 python ; done  
exit
```

# Summary

- **Most ML/DL packages on Comet are enabled via Singularity**
- **Single node runs are easy to wrap with “singularity exec” commands**
- **Can launch and use Jupyter notebooks from within Singularity containers**
- **Multi-node runs need some work but feasible in most cases.**